

# Introducing the Tidyverse

Matteo Fasiolo

University of Bristol

# Overview

## What is the “Tidyverse”?

The Tidyverse is a set of inter-compatible R packages sharing the same programming philosophy.

## What does it offer?

A coherent and consistent R programming framework for:

- 1 readr: loading data into R
- 2 tidyr: putting your data in a tidy format
- 3 dplyr: manipulating data
- 4 ggplot2: building data visualizations
- 5 purrr: doing functional programming in R

As well as for more specific purposes:

- 1 stringr (strings), forcats (factor variables), lubridate (dates and times), ...

# Overview

You will probably never use the whole Tidyverse.

Here we cover the basics:

- 1 `magrittr`: how pipes `%>%` enhance code readability
- 2 `ggplot2`: how layered graphics work
- 3 `dplyr`: how to modify dataframes in a clear and consistent manner
- 4 `tidyr`: how to reshape your data for plotting and modelling

Why do you need to know this:

- 1 by manipulating **real data** in base R (e.g., `tapply`, `sapply`, ...) routinely, you might go crazy
- 2 industry jobs
- 3 Tidyverse packages illustrates many programming concepts that are widely useful

# Online material 1: pipes

The pipe operator `%>%` is provided by `magrittr`.

Consider plotting  $\sqrt{|\cos(x)|}$  on a grid:

```
x <- seq(0, 2*pi, by = 0.01)

plot( sqrt ( abs( cos(x) ) ) )
```

The piped equivalent:

```
x %>% cos %>% abs %>% sqrt %>% plot
```

How does this work? Consider `f(a1, a2, a3)`:

```
x %>% f()           equivalent to   f(a1 = x)

x %>% f(3, 5)       equivalent to   f(a1 = x, a2 = 3, a3 = 5)
```

# Online material 1: pipes

How is this useful? Electricity demand example:

```
data(UKload)

plot(NetDemand ~ Posan,
     transform(head(subset(UKload, Dow == "lundi",
                           select = c("NetDemand", "Posan")), 100),
              Posan = Posan * 365))
```

The piped equivalent:

```
UKload %>%
  subset(Dow == "lundi",
        select = c("NetDemand", "Posan")) %>%
  head(100) %>%
  transform(Posan = Posan * 365) %>%
  plot(NetDemand ~ Posan, data = .)
```

# Online material 1: pipes

The advantages are:

- aesthetic (subjective)
- improved clarity → fewer errors
- compatibility with rest of Tidyverse

The material covers also other pipes: %<>%, %\$%, ...

**Note** that there is a danger of going to far, e.g.:

```
sqrt( abs( cos(x) ) )
```

is closer to  $\sqrt{|\cos(x)|}$  than

```
x %>% cos %>% abs %>% sqrt
```

## Online material 2: ggplot2

We briefly introduce ggplot2:

- 1 how to build basic plots
- 2 how to add layers and facets

A scatterplot in base R:

```
data(mcycle)

plot(x = mcycle$times, y = mcycle$accel)
```

Function plot is called for its side effects:

```
tmp <- plot(x = mcycle$times, y = mcycle$accel)

tmp

## NULL
```

## Online material 2: ggplot2

A scatterplot in ggplot2:

```
# Building object
pl <- ggplot(data = mcycle)

# Adding layers
pl <- pl + geom_point(mapping = aes(x=times, y=accel))

# Rendering on screen
pl
```

Basic template:

```
ggplot(data = <data.frame>) +  
  <geom_layer>(mapping = aes(<variables_map>))
```



## Online material 3: ggplot2 case study

Ok, ggplot2 plots are pretty, but what else?

mgcViz case study provides some motivation.

Context is GAM modelling with mgcv:

```
fitG <- gam(Demand ~ Dow + s(Posan) + s(wM) + s(Trend),  
            data = UKload)
```

We plot the effect of Posan by doing:

```
plot(fitG, select = 1)
```

This will call `plot.gam`.

## Online material 3: ggplot2 case study

plot.gam does its job but:

```
args(plot.gam)
```


```
function (x, residuals = FALSE, rug = NULL, se = TRUE,
  pages = 0, select = NULL, scale = -1, n = 100,
  n2 = 40, n3 = 3, pers = FALSE, theta = 30,
  phi = 30, jit = FALSE, xlab = NULL, ylab = NULL,
  main = NULL, ylim = NULL, xlim = NULL,
  too.far = 0.1, all.terms = FALSE, shade = FALSE,
  shade.col = "gray80", shift = 0, trans = I,
  seWithMean = FALSE, unconditional = FALSE,
  by.resids = FALSE, scheme = 0, ...)
```

Quite a lot of arguments and...

# Online material 3: ggplot2 case study

Quite a lot of arguments and:

- difficult to add new features
- cannot control properties of elements
- order in which elements are rendered is fixed



## Online material 3: ggplot2 case study

Quite a lot of arguments and:

- difficult to add new features
- cannot control properties of elements
- order in which elements are rendered is fixed

mgcViz wraps GAM object allows us to do:

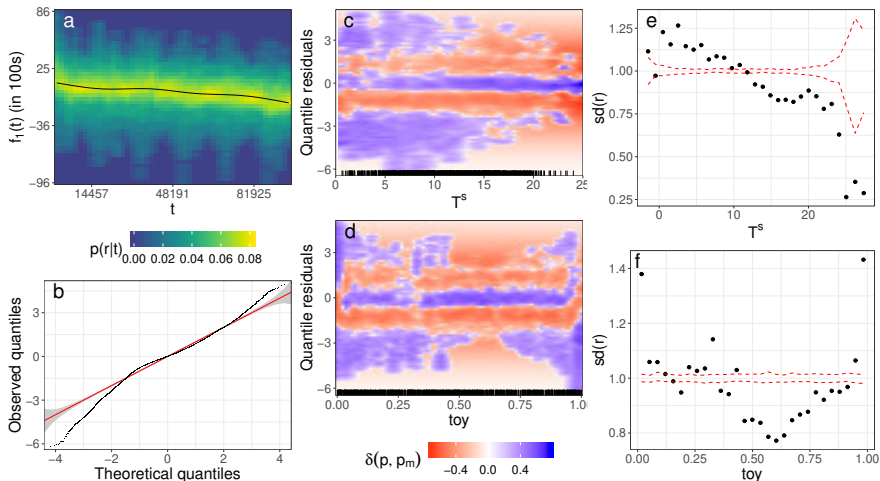
```
plot(fitG, 1) + l_fitLine(col = "red") +  
               l_ciLine(linetype = 2) +  
               xlim(0.25, 0.75)
```

where, e.g., `l_fitLine` is a wrapper for `geom_line`.

Amazingly, this solves all the problems above.

# Online material 3: ggplot2 case study

Personal conclusion is that, if you want to build an flexible and extensible graphical library in R, ggplot2 might be the way to go.



## Online material 4: dplyr

ggplot2 wants you to provide a `data.frame`:

```
ggplot(data = <data.frame>) +  
  <geom_layer>(mapping = aes(<variables_map>))
```

This is annoying when you just want `plot(x, y)`, `hist(x)`...

For more complex plots, effort is justified.

Many modelling functions (e.g. `lm`, `glm`, `gam`) have same requirement.

`dplyr` and `tidyr` help us building the “right” `data.frame` for visualization and modelling.

## Online material 4: dplyr

The online notes focus on the basics. Why? Well...

```
length( getNamespaceExports("dplyr") ) # Oct 2022  
287
```

UK demand example:

```
head(UKload)  
      NetDemand    wM wM_s95 Posan      Dow      Trend  
25      38353 6.05    5.56  0.00   samedi 1293879600  
73      41192 2.80    3.23  0.00 dimanche 1293966000  
121     43442 2.10    1.86  0.01   lundi  1294052400
```

## Online material 4: dplyr

Example of dplyr code:

```
UKload %>% select(NetDemand, wM, Dow, Posan) %>%  
  filter(wM < 5 & Dow == "lundi") %>%  
  arrange(desc(wM)) %>%  
  ggplot() +  
  geom_point(aes(x=wM, NetDemand))
```

Base R equivalent:

```
d0 <- UKload[ , c("NetDemand", "wM", "Dow", "Posan")]  
d0 <- d0[d0$wM < 5 & d0$Dow == "lundi", ]  
d0 <- d0[rev(order(d0$wM)), ]  
plot(NetDemand ~ wM, d0)
```



## Online material 4: dplyr

A more interesting example:

```
d0 <- UKload %>% mutate(wk = week(Date)) %>%  
  group_by(Year, wk) %>%  
  summarise(TotDemand = sum(NetDemand),  
            tempMax = max(wM),  
            tempMin = min(wM),  
            nHoly = sum(Holy=="1"))
```

Base R equivalent:

```
d0 <- UKload  
d0$wk <- week(UKload$Date)  
???
```

There surely is a good base R solution but dplyr code generally clearer and more concise.

# Online material 5: data reshaping with tidyr and dplyr

We consider Irish electricity smart meter data:

```
data(Irish)
indCons <- Irish$indCons
head(indCons)
##           I1002 I1003 I1004 I1005 I1013 I1015 I1018 ...
## 8114 0.022 0.593 2.002 0.755 0.035 0.398 0.547 ...
## 8115 0.133 0.707 1.602 0.898 0.112 0.689 0.603 ...
## 8116 0.094 0.684 1.525 0.736 0.046 0.407 0.511 ...
```

But we need data in “long” format for ggplotting and modelling:

```
##           ID    dem
## 1 I1002 0.022
## 2 I1002 0.133
## 3 I1002 0.094
## 4 I1002 0.023
```

## Online material 5: data reshaping with tidyr and dplyr

Easily done with `tidyr::pivot_longer`:

```
longDat <- indCons %>% pivot_longer(cols = everything(),
  names_to = "ID", values_to = "dem") %>% arrange(ID)
head(longDat)
##           ID    dem
## 1 I1002 0.022
## 2 I1002 0.133
## 3 I1002 0.094
```

But there is a memory price to pay:

```
indCons %>% object.size %>% format(units = "MB")
# "12.9 Mb"

longDat %>% object.size %>% format(units = "MB")
# "25.6 Mb"
```

## Online material 5: data reshaping with tidyr and dplyr

Opposite transformation achieved by `tidyr::pivot_wider`:

```
wideDat <- longDat %>% pivot_wider(names_from = "ID",  
                                   values_from = "dem")  
  
head(wideDat)  
  
# A tibble: 16,799 x 101  
  I1002 I1003 I1004 I1005 I1013 I1015 I1018  
  <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>  
1 0.022 0.593  2.00 0.755 0.035 0.398 0.547  
2 0.133 0.707  1.60 0.898 0.112 0.689 0.603  
3 0.094 0.684  1.52 0.736 0.046 0.407 0.511  
4 0.023 0.563  1.39 0.738 0.036 0.223 0.593  
5 0.133 0.489  1.22 0.849 0.065 0.132 0.570  
# with 1.679e+04 more rows, and 94 more variables
```

## Online material 5: data reshaping with tidyr and dplyr

Another common task is merging data frames, e.g.:

```
extra <- as_tibble( Irish$extra )  
head(extra)  
## # A tibble: 6 x 6  
##   time    toy dow   holy    tod  temp  
##   <int> <dbl> <fct> <lgl> <dbl> <dbl>  
## 1     1  0.986 Wed   FALSE     0     4  
## 2     2  0.986 Wed   FALSE     1     4  
## 3     3  0.986 Wed   FALSE     2     4  
## 4     4  0.986 Wed   FALSE     3     4  
## 5     5  0.986 Wed   FALSE     4     4  
## 6     6  0.986 Wed   FALSE     5     4
```

## Online material 5: data reshaping with tidyr and dplyr

Here merging is easy:

```
allDat <- cbind(longDat, extra)
head(allDat)
```

##	ID	dem	time	toy	dow	holy	tod	temp
## 1	I1002	0.022	1	0.9863014	Wed	FALSE	0	4
## 2	I1002	0.133	2	0.9863014	Wed	FALSE	1	4
## 3	I1002	0.094	3	0.9863014	Wed	FALSE	2	4
## 4	I1002	0.023	4	0.9863014	Wed	FALSE	3	4
## 5	I1002	0.133	5	0.9863014	Wed	FALSE	4	4
## 6	I1002	0.090	6	0.9863014	Wed	FALSE	5	4

But remember about memory costs ...

## Online material 5: data reshaping with tidyr and dplyr

But how to add also the customer survey data:

```
survey <- as_tibble( Irish$survey )  
head(survey)
```

```
# A tibble: 6 x 12
```

	ID	meanDem	SCLASS	OWNERSHIP	YEAR	HEAT.HOME	HEAT.WAT
	<chr>	<dbl>	<fct>	<fct>	<dbl>	<fct>	<fct>
1	I1002	0.208	DE	0	1975	Other	Elec
2	I1003	0.622	C1	0	2004	Other	Other
3	I1004	0.962	C1	0	1987	Other	Elec
4	I1005	0.640	C1	0	1930	Other	Other
5	I1013	0.241	C2	0	2003	Other	Elec
6	I1015	0.463	DE	R	1989	Elec	Other

```
# with 5 more variables: WINDOWS.doubleglazed <fct>,  
# HOME.APPLIANCE..White.goods. <dbl>, Code <int>,
```

## Online material 5: data reshaping with tidyr and dplyr

Solution is offered by `dplyr::left_join`:

```
allDat <- allDat %>% left_join(survey, by = "ID") %>%  
  as_tibble()  
  
## # A tibble: 6 x 20  
##   ID      dem    time    toy dow   holy    tod  SCLASS  
##   <chr> <dbl>   <int> <dbl> <fct> <lgl> <dbl> <FCT>  
## 1 I1002 0.022     1 0.986 Wed   FALSE     0    DE  
## 2 I1002 0.133     2 0.986 Wed   FALSE     1    DE  
## 3 I1002 0.094     3 0.986 Wed   FALSE     2    DE  
## 4 I1002 0.023     4 0.986 Wed   FALSE     3    DE  
## 5 I1002 0.133     5 0.986 Wed   FALSE     4    DE  
## 6 I1002 0.09      6 0.986 Wed   FALSE     5    DE  
## #   with 10 more variables: OWNERSHIP <fct>,  
## #   BUILT.YEAR <dbl>, HEAT.HOME <fct>, HEAT.WAT <fct>,  
## #   WINDOWS.doubleglazed <fct>
```



## Online material 5: data reshaping with tidyr and dplyr

Now we have all the info in one data.frame but:

```
indCons %>% object.size() %>% format("MB")  
## [1] "12.9 Mb"
```

```
survey %>% object.size() %>% format("MB")  
## [1] "0.3 Mb"
```

```
extra %>% object.size() %>% format("MB")  
## [1] "0.6 Mb"
```

```
allDat %>% object.size() %>% format("MB")  
## [1] "217.9 Mb"
```

## Further topics

Online material covers basics, for details see “R for Data Science” book.  
Available online at <https://r4ds.had.co.nz/>.

For a skeptical point of view on the Tidyverse, see:

<https://github.com/matloff/TidyverseSkeptic>

- Tidyverse “bad” when teaching to ppl with no background
- Tidyverse averse to \$, `[[ ]]`, loops and `plot()`
- Tidyverse is advertised by Rstudio (data.table example)
- too many functions: `mutate`, `mutate_`, `mutate_all`, `mutate_at`, `mutate_each`, `mutate_each_`, `mutate_if`, `transmute`, `transmute_`, `transmute_all`, `transmute_at`

Probably the most useful thing to do for you is to experiment the Tidyverse packages on real data.

The data sets used in the online notes are in `qgam` and `electBook` packages.

An ideal data set is also provided in this Kaggle challenge:

<https://www.kaggle.com/c/ashrae-energy-prediction>

# Computer lab

Consumption data:

```
head(train)
  building_id meter      timestamp meter_reading
1           0     0 2016-01-01 00:00:00         0
2           1     0 2016-01-01 00:00:00         0
3           2     0 2016-01-01 00:00:00         0
4           3     0 2016-01-01 00:00:00         0
```

Building information:

```
head(buildMeta)
  site_id building_id primary_use square_feet year_built
1         0           0  Education       7432      2008
2         0           1  Education       2720      2004
3         0           2  Education       5376      1991
```

# Computer lab

Weather data:

```
head(weather)
  site_id air_temperature cloud_coverage wind_speed
1      0             25.0              6         0.0
2      0             24.4             NA         1.5
3      0             22.8              2         0.0
```

The data could get big:

```
train %>% object.size %>% format("MB")
"386.4 Mb"
weather %>% object.size %>% format("MB")
"9.3 Mb"

nlevels(train$building_id)
1449
```

# References I

Hastie, T. and R. Tibshirani (1990). *Generalized Additive Models*, Volume 43. CRC Press.