

Portfolio 3

Sam Bowyer

2022-10-13

Vectorisation

R is an interpreted language, so loops (particularly nested loops) can be very slow. Instead, it is often useful to use built-in vector operations, which tend to be written in C/C++ and hence are compiled and run much faster.

As example consider the two functions below that both compute the minimum value that the sign function takes from elements in a vector:

```
minsin1 <- function(x){  
  m = Inf  
  for (i in 1:length(x)){  
    if (sin(x[i]) < m){}  
    m = sin(x[i])  
  }  
  return (m)  
}  
  
minsin2 <-function(x) min(sin(x))
```

Then running these two functions on the same input, we find that the second, vectorised version runs much faster:

```
x = 1:1e7  
system.time(minsin1(x))  
  
##      user  system elapsed  
##  0.668    0.008    0.676  
  
system.time(minsin2(x))  
  
##      user  system elapsed  
##  0.120    0.028    0.149
```

Vectorisation is then particularly useful when we're working with matrices, or even higher dimensional arrays—more dimensions typically mean more nested loops in unvectorised code, greatly slowing things down. Consider the two functions below which sum the rows of a matrix before applying the previous `minsin` functions to the resulting vector.

```
minsin_matrix1 <- function(X){  
  m = Inf  
  for (i in 1:nrow(X)){  
    total = 0  
    for (j in 1:ncol(X)){  
      total = total + X[i,j]  
    }  
  }  
}
```

```

    if (sin(total) < m){
      m = sin(total)
    }
  }
  return (m)
}

minsin_matrix2 <- function(X) min(sin(rowSums(X)))

```

Again, by running these functions on the same data we find that vectorisation has drastically sped up execution.

```

X = matrix(1:1e8, 1e3, 1e5)
system.time(minsin_matrix1(X))

```

```

##      user  system elapsed
##    2.684    0.000    2.685

system.time(minsin_matrix2(X))

```

```

##      user  system elapsed
##    0.209    0.000    0.208

```

Useful Functions

R includes a variety of functions which help us perform operations on vectors.

apply

`apply(X, MARGIN, FUN, ...)`: applies the function `FUN` to an array of dimension 2 or more using the dimensions given in the list `MARGIN` (in which 1 represents rows, 2 represents columns, `c(1,2)` represents both, etc.).

```

x <- matrix(1:12, 3, 4)
print(x)

```

```

##      [,1] [,2] [,3] [,4]
## [1,]    1    4    7   10
## [2,]    2    5    8   11
## [3,]    3    6    9   12

```

```

apply(x, c(1,2), minsin2)

```

```

##      [,1]      [,2]      [,3]      [,4]
## [1,] 0.8414710 -0.7568025 0.6569866 -0.5440211
## [2,] 0.9092974 -0.9589243 0.9893582 -0.9999902
## [3,] 0.1411200 -0.2794155 0.4121185 -0.5365729

```

```

apply(x, 1, minsin2)

```

```

## [1] -0.7568025 -0.9999902 -0.5365729

```

```

apply(x, 2, minsin2)

```

```

## [1] 0.1411200 -0.9589243 0.4121185 -0.9999902

```

For an example with a 3-dimensional array:

```
x <- array(1:12, c(2, 3, 2))
print(x)
```

```
## , , 1
##
##      [,1] [,2] [,3]
## [1,]    1    3    5
## [2,]    2    4    6
##
## , , 2
##
##      [,1] [,2] [,3]
## [1,]    7    9   11
## [2,]    8   10   12
```

```
apply(x, 3, sum)
```

```
## [1] 21 57
```

```
apply(x, c(1,2), sum)
```

```
##      [,1] [,2] [,3]
## [1,]    8   12   16
## [2,]   10   14   18
```

```
apply(x, c(2,3), sum)
```

```
##      [,1] [,2]
## [1,]    3   15
## [2,]    7   19
## [3,]   11   23
```

```
apply(x, c(1,3), sum)
```

```
##      [,1] [,2]
## [1,]    9   27
## [2,]   12   30
```

lapply

lapply(X, FUN, ...): works like `apply` but can be used on vectors and lists, and also returns a list.

```
lapply(1:4, sqrt)
```

```
## [[1]]
## [1] 1
##
## [[2]]
## [1] 1.414214
##
## [[3]]
## [1] 1.732051
##
## [[4]]
## [1] 2
```

```
x <- list(a=1:3, b=c(TRUE, TRUE, FALSE), c=2:-1)
lapply(x, minsin2)
```

```
## $a
## [1] 0.14112
##
## $b
## [1] 0
##
## $c
## [1] -0.841471
```

Note that in the following code execution we find that `lapply` is slower than both vectorised code and a `for` loop.

```
func <- function(x) sqrt(x^2)
func_lapply <- function(x) lapply(x, func)
func_loop <- function(x){
  out = rep(NA, length(x))
  for (i in seq_len(length(x))){
    out[i] = func(x[i])
  }
  return(out)
}
x = 1:1e7
```

```
system.time(func(x))
```

```
##      user  system elapsed
##  0.033   0.016   0.049
```

```
system.time(func_lapply(x))
```

```
##      user  system elapsed
##  5.524   0.156   5.680
```

```
system.time(func_loop(x))
```

```
##      user  system elapsed
##  2.449   0.000   2.449
```

sapply

`sapply(X, FUN, ...)`: works like `sapply` but simplifies the output before returning.

```
sapply(1:4, sqrt)
```

```
## [1] 1.000000 1.414214 1.732051 2.000000
```

```
x <- list(a=1:3, b=c(TRUE, TRUE, FALSE), c=2:-1)
sapply(x, minsin2)
```

```
##      a      b      c
## 0.141120 0.000000 -0.841471
```

mapply

`mapply(FUN, ...)`: the (potentially multiple) arguments given as `...` are used to run the function `FUN`.

```
mapply(sqrt, 1:4)
```

```
## [1] 1.000000 1.414214 1.732051 2.000000
```

```
mapply(function(x,y,z) x * y + z, c(1, 10, 100), c(2,3,4), c(0, 1, 2))
```

```
## [1] 2 31 402
```

Map

This works very similarly to `mapply`.

```
Map(rep, 1:3, 4:6)
```

```
## [[1]]  
## [1] 1 1 1 1  
##  
## [[2]]  
## [1] 2 2 2 2  
##  
## [[3]]  
## [1] 3 3 3 3 3
```

Though it is ever so slightly faster.

```
system.time(Map(func, 1:1e7))
```

```
##      user  system elapsed  
##  5.313    0.000   5.312
```

```
system.time(mapply(func, 1:1e7))
```

```
##      user  system elapsed  
##  5.840    0.091   5.934
```

Reduce

`Reduce(FUN, X)` applies `FUN` to consecutive pairs of elements in a vector iteratively until a single element is left.

```
Reduce(rep, 1:3)
```

```
## [1] 1 1 1 1 1 1
```

Above, `Reduce` has first executed `rep(1,2)` to obtain the vector `c(1,1)` and has then executed `rep(c(1,1), 3)` to obtain the output.

Below, `Reduce` is used to write the elements of a list as the digits in a number.

```
Reduce(function(a,b) 10*a + b, 1:6)
```

```
## [1] 123456
```

Filter

`Filter(FUN, X)` removes any elements from the vector `X` who do not evaluate to `TRUE` under the function `FUN`.

```
Filter(function(x) sqrt(x) %% 1 == 0, 1:30)
```

```
## [1] 1 4 9 16 25
```

```
Filter(function(x) sqrt(x^2) == x, -10:10)
```

```
## [1] 0 1 2 3 4 5 6 7 8 9 10
```

Parallel Programming