

# Portfolio 2 (Extended) - Integrating R and C++

Sam Bowyer

2023-05-29

Typically, interfacing R with C++ is done through Rcpp. In the final section of this extended portfolio we'll use Rcpp, however, first we'll look into how we can use the raw C API given by R to give us a better sense of what Rcpp is doing.

## Simulation-based inference on the Ricker model

First we will be considering inference on the Ricker model, a simple model for population dynamics:

$$y_{t+1} = ry_t e^{-y_t}$$

where  $y_t > 0$  is the population size at time  $t$  and  $r > 0$  is the growth rate.

Below we present an implementation of the Ricker model in R, for an initial population size `y0` running for `n` time steps after an initial `nburn` time steps which are discarded.

```
rickerSimulR <- function(n, nburn, r, y0 = 1){
  y <- numeric(n)
  yx <- y0

  # Burn in phase
  if(nburn > 0){
    for(ii in 1:nburn){
      yx <- r * yx * exp(-yx)
    }
  }

  # Simulating and storing
  for(ii in 1:n){
    yx <- r * yx * exp(-yx)
    y[ii] <- yx
  }

  return(y)
}
```

## Question 1

We can write a version of the above model in C as follows:

```
cat ./rickerSimul.c
```

```
## #include <R.h>
## #include <Rinternals.h>
## #include <Rmath.h>
##
```

```

## SEXP rickerSimul(SEXP num, SEXP numburn, SEXP rate, SEXP initialPop){
##     double *xys;
##     int n, nburn;
##     double r, y0;
##     SEXP ys;
##
##     n = INTEGER(num)[0];
##     ys = PROTECT(allocVector(REALSXP, n));
##     xys = REAL(ys);
##
##     nburn = INTEGER(numburn)[0];
##     r = REAL(rate)[0];
##     y0 = REAL(initialPop)[0];
##
##     double yx = y0;
##
##     // Burn in phase
##     if(nburn > 0){
##         for(int i = 0; i < nburn; i++){
##             yx = r * yx * exp(-yx);
##         }
##     }
##
##     // Simulating and storing
##     for(int i=1; i < n; i++){
##         yx = r * yx * exp(-yx);
##         xys[i] = yx;
##     }
##
##     UNPROTECT(1);
##
##     return ys;
## }

```

Then we compile it with in R with the following line

```
system("R CMD SHLIB rickerSimul.c")
```

This has created two files, a .o and a .so file.

```
ls rickerSimul.*
```

```

## rickerSimul.c
## rickerSimul.o
## rickerSimul.so

```

We then load the .so file into R and call it using .Call

```

dyn.load("rickerSimul.so")
is.loaded("rickerSimul")

```

```
## [1] TRUE
```

```

n = 25L
nburn=5L
r = 5
y0 = 4
c_output = .Call("rickerSimul", n, nburn, r, y0)

```

```
c_output
```

```
## [1] 4.658982e-310 1.569201e+00 1.633629e+00 1.594584e+00 1.618446e+00
## [6] 1.603932e+00 1.612787e+00 1.607394e+00 1.610682e+00 1.608679e+00
## [11] 1.609900e+00 1.609156e+00 1.609610e+00 1.609333e+00 1.609502e+00
## [16] 1.609399e+00 1.609462e+00 1.609423e+00 1.609447e+00 1.609433e+00
## [21] 1.609441e+00 1.609436e+00 1.609439e+00 1.609437e+00 1.609438e+00
```

And importantly we can see that this has produced the same results as the R version:

```
r_output = rickerSimulR(n, nburn, r, y0)
max(abs(c_output - r_output))
```

```
## [1] 1.569201
```

But importantly, the C version is much faster than the R version.

```
rickerSimulC_ <- function() .Call("rickerSimul", 100L, 20L, 10, 1)
rickerSimulR_ <- function() rickerSimulR(100L, 20L, 10, 1)
```

```
library(microbenchmark)
microbenchmark(rickerSimulC_(), rickerSimulR_(), times=10000)
```

```
## Unit: microseconds
##      expr    min      lq     mean  median      uq      max neval
## rickerSimulC_() 1.276 1.322 1.685081   1.374 1.6155 552.512 10000
## rickerSimulR_() 5.123 5.442 6.009669   5.590 6.2580 352.197 10000
```

## Question 2

Now suppose we have noisy observations from the Ricker model:


$$z_t = y_t e^{\epsilon_t} \text{ where } \epsilon_t \sim N(0, \sigma^2)$$

```
nburn <- 100L
n <- 50L

y0_true <- 1
sig_true <- 0.1
r_true <- 10

Ntrue <- rickerSimulR(n = n, nburn = nburn, r = r_true, y0 = y0_true)
yobs <- Ntrue * exp(rnorm(n, 0, sig_true))

plot(yobs, type = 'b')
```



We then write the following function in C to calculate the log likelihood of the data (this function is in the file `rickerLLK.c`):

```
cat rickerLLK.c
```

```
## #include <R.h>
## #include <Rinternals.h>
## #include <Rmath.h>
##
## SEXP rickerLLK(SEXP observed, SEXP simulated, SEXP sigma){
##     double *yobs, *ysim, sig, *lik;
##     int n;
##
##     SEXP LLK;
##
##     yobs = REAL(observed);
##     ysim = REAL(simulated);
##     sig = REAL(sigma)[0];
##     n = length(observed);
##
##     LLK = PROTECT(allocVector(REALSXP, 1));
##
##     lik = REAL(LLK);
##     lik[0] = 0;
##
##     for (int i = 1; i < n; i++){
##         lik[0] = lik[0] + dnorm(log(yobs[i]/ysim[i]), 0, sig, 1);
##     }
##
##     UNPROTECT(1);
##     return LLK;
## }
```

```
system("R CMD SHLIB rickerLLK.c")
dyn.load("rickerLLK.so")
is.loaded("rickerLLK")
```

```
## [1] TRUE
```

Next we wrap the likelihood calculation in an R function that takes in the logarithm of  $r$ ,  $\sigma$  and  $y_0$  as well as  $y_{\text{obs}}$  and  $n_{\text{burn}}$ :

```
myLikR <- function(logr, logsig, logy0, yobs, nburn){
  n <- length(yobs)
  r <- exp(logr)
  sig <- exp(logsig)
  y0 <- exp(logy0)

  ysim <- .Call("rickerSimul", n, nburn, r, y0)

  llk <- .Call("rickerLLK", yobs, ysim, sig)


  return( llk )
}

myLikR(log(r_true), log(sig_true), log(y0_true), yobs, nburn)
```

```
## [1] -4482.756
```

With this, we can then sample from the posterior distribution of  $\log(r)$ ,  $\log(\sigma)$ ,  $\log(y_0)$  by providing `myLikR` to a Metropolis-Hastings algorithm (using the `metrop` function):

```
library(mcmc)
par(mfrow=c(1,3))
samples = metrop(function(params) myLikR(params[1], params[2], params[3], yobs, nburn),
  initial = c(1,1,1),
  nbatch = 10000,
  scale = 0.05)
plot(density(samples$batch[,1]), main="log r Posterior")
plot(density(samples$batch[,2]), main="log sig Posterior")
plot(density(samples$batch[,3]), main="log y0 Posterior")
```



Recall the true log values of the parameters:

```
log(r_true); log(sig_true); log(y0_true)
```

```
## [1] 2.302585
## [1] -2.302585
## [1] 0
```

We see that we've got sharp posteriors around the correct values for **r** and **sig**, however, a much broader posterior for **y0** that isn't very close to the true value of 0. This is because although **r** and **sig** have a great impact on every observation, as the simulation goes on for many samples, the impact of the initial state, **y0**, drastically decreases, meaning it is harder to infer from the data.


### Question 3

Now assume that the data has been simulated as follows:

```
r_true <- 44

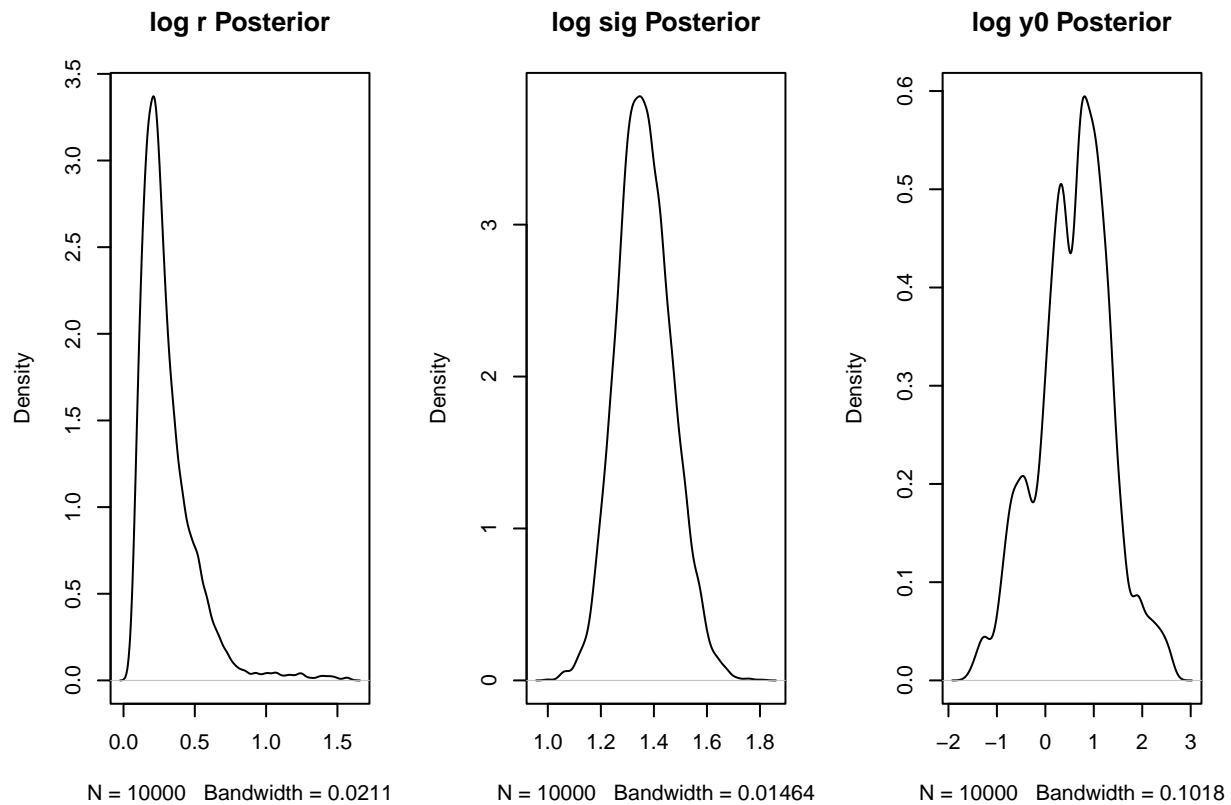
Ntrue <- rickerSimulR(n = n, nburn = nburn, r = r_true, y0 = y0_true)
yobs <- Ntrue * exp(rnorm(n, 0, sig_true))

plot(yobs, type = 'b')
```



Attempting to run the MH algorithm as before, we notice that the chain is failing to mix properly, resulting in broader, less accurate posteriors:

```
par(mfrow=c(1,3))
samples = metrop(function(params) myLikR(params[1], params[2], params[3], yobs, nburn),
  initial = c(1,1,1),
  nbatch = 10000,
  scale = 0.05)
plot(density(samples$batch[,1]), main="log r Posterior")
plot(density(samples$batch[,2]), main="log sig Posterior")
plot(density(samples$batch[,3]), main="log y0 Posterior")
```



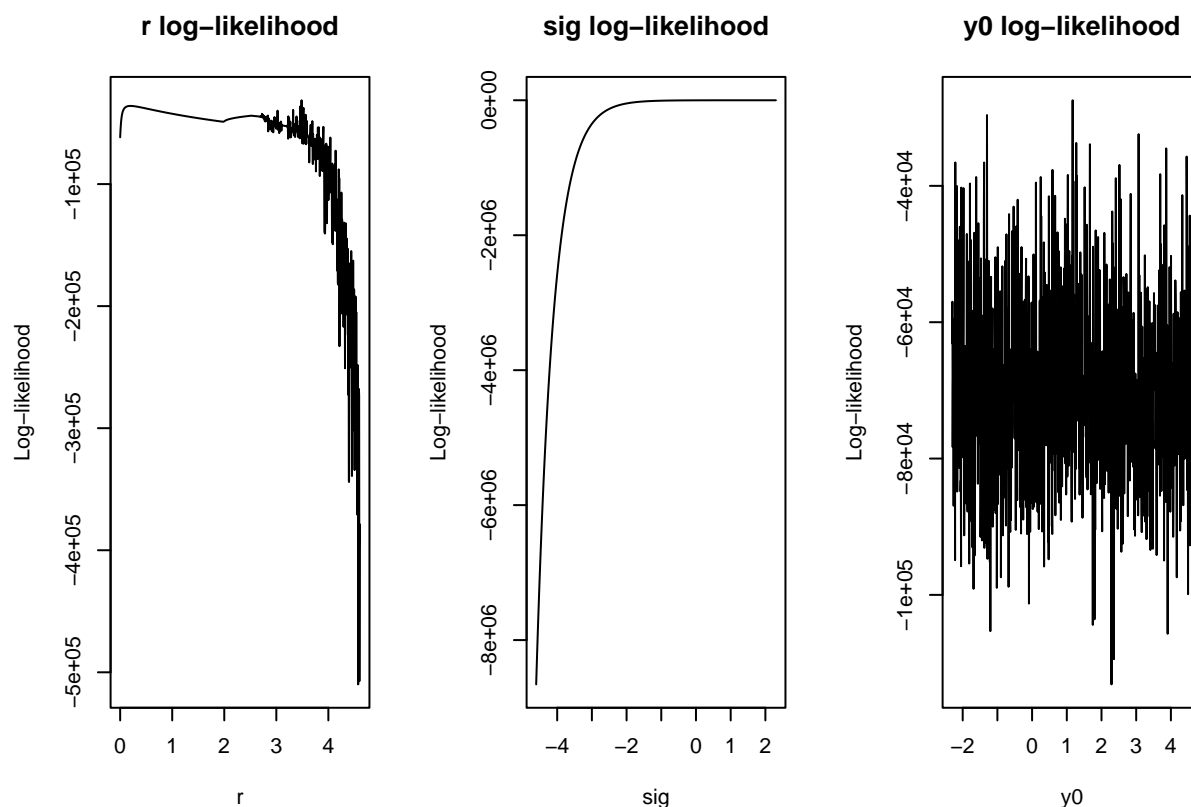
To investigate this further, we'll look at the slices of likelihood with respect to each parameter (keeping the other two parameters fixed at a time to their true values).

```
par(mfrow=c(1,3))

r_seq = seq(log(1),log(100), length.out=1000)
sig_seq = seq(log(0.01),log(10), length.out=1000)
y0_seq = seq(log(0.1),log(100), length.out=1000)

plot(r_seq, sapply(r_seq, function(x) myLikR(x, log(sig_true), log(y0_true), yobs, nburn)),
     main="r log-likelihood", xlab="r", ylab="Log-likelihood", type='l')
plot(sig_seq, sapply(sig_seq, function(x) myLikR(log(r_true), x, log(y0_true), yobs, nburn)),
     main="sig log-likelihood", xlab="sig", ylab="Log-likelihood", type='l')
plot(y0_seq, sapply(y0_seq, function(x) myLikR(log(r_true), log(sig_true), x, yobs, nburn)),
     main="y0 log-likelihood", xlab="y0", ylab="Log-likelihood", type='l')
```





We see above that the log-likelihood for  $r$  gets very unpleasant for  $r > \exp(3)$ , meaning that the MH algorithm isn't able to mix well with the chaotic behaviour of the model for small changes in  $r$  and  $y_0$ .

## Question 4

To deal with this, we will now assume that we know the true value of  $\sigma$  and that  $y_0 \sim \text{Unif}(1, 10)$  (but in particular we don't care about the value of  $y_0$ ). We can come up with another way to express the likelihood of our data by assuming that the sample mean  $s_1$  and standard deviation  $s_2$  are independently normally distributed:

$$s_1 \sim \mathcal{N}(\mu_1, \tau_1^2), \quad s_2 \sim \mathcal{N}(\mu_2, \tau_2^2).$$

Then the desired likelihood  $p(s_1, s_2 | r)$  is simply the product of these two normal densities, where  $\mu_1, \mu_2, \tau_1, \tau_2$  are functions of  $r$ . We can sample the corresponding posterior by simulation (`nsim` times for a given value of `logr`) via the following function:

```
synllk <- function(logr, nsim){

  r <- exp(logr)
  s1 <- s2 <- numeric(nsim)
  y0 <- runif(nsim, 1, 10)

  # Note: sigma is assumed to be known!
  for(ii in 1:nsim){
    ysim <- rickerSimulR(n = n, nburn = nburn, r = r, y0 = y0[ii]) * exp(rnorm(n, 0, sig_true))
    s1[ii] <- mean(ysim)
    s2[ii] <- sd(ysim)
  }

  out <- dnorm(mean(yobs), mean(s1), sd(s1), log = TRUE) +
```

```

    dnorm(sd(yobs), mean(s2), sd(s2), log = TRUE)

    return( out )
}

```


Then using the same Metropolis-Hastings approach as before we can sampling from an estimate of the posterior:

```

samples = metrop(function(r) synllk(r, 100), initial=log(40), nbatch=10000)
plot(density(exp(samples$batch[,1])), main="r Approximate Posterior")

```

## r Approximate Posterior



N = 10000 Bandwidth = 0.5093

This does a much better job at estimating the true value of  $r = 44$  (though this being an approximation of the posterior, it is somewhat broad and multimodal). However, the implementation of this synthetic likelihood in R is very slow, and so below we provide an implementation of `synllk` and `rickerSimul` using Rcpp, stored in a file `rickerRcpp.cpp`.

```
cat rickerRcpp.cpp
```

```

## #include <Rcpp.h>
## using namespace Rcpp;
##
## // [[Rcpp::export]]
## NumericVector rickerSimul_Rcpp(const int n, const int nburn, const double r, const double y0){
##     // vector to return
##     NumericVector y(n);
##
##     double yx = y0;
##
##     // burn-in
##     if(nburn > 0){
##         for(int i=0; i<=nburn; i++){

```

```

##         yx = r * yx * exp(-yx);
##     }
## }
##
## // run simulation and store values
## y[0] = yx;
## for(int i=1; i<n; i++){
##     yx = r * yx * exp(-yx);
##     y[i] = yx;
## }
##
## return y;
## }
##
## // [[Rcpp::export]]
## NumericVector synllk_Rcpp(const double logr, const int nsim, const NumericVector yobs){
##     NumericVector ysim;
##     NumericVector s1(nsim), s2(nsim);
##     NumericVector y0(nsim);
##     NumericVector out;
##
##     double r = exp(logr);
##
##     y0 = runif(nsim, 0, 10);
##     for(int i=0; i<nsim; i++){
##         // assume we know sd = 0.1
##         ysim = rickerSimul_Rcpp(50, 100, r, y0[i]) * exp(rnorm(50, 0, 0.1));
##         s1[i] = mean(ysim);
##         s2[i] = sd(ysim);
##     }
##
##     // add log-likelihoods of our two summary statistics
##     out = R::dnorm(mean(yobs), mean(s1), sd(s1), true) + R::dnorm(sd(yobs), mean(s2), sd(s2), true);
##
##     return out;
## }

```


Then we can load this in R with the `sourceCpp` function:

```

library(Rcpp)
sourceCpp("rickerRcpp.cpp")
samples = metrop(function(r) synllk_Rcpp(r, 100, yobs), initial=log(40), nbatch=10000)
plot(density(exp(samples$batch[,1])), main="r Approximate Posterior")

```

## r Approximate Posterior



As we can see, this produces similar results to before and, importantly, the Rcpp version of `synllk` is much faster than the pure R version:

```
synllk_Rcpp_ <- function() synllk_Rcpp(r_true, 100, yobs)
synllk_R_ <- function() synllk(r_true, 100)
microbenchmark(synllk_Rcpp_(), synllk_R_(), times=10000)
```

```
## Unit: microseconds
##      expr      min       lq      mean    median      uq      max neval
## synllk_Rcpp_() 186.476 200.056 222.4572 207.9165 227.5625 2832.971 10000
##   synllk_R_() 1453.666 1591.685 1751.5768 1640.6925 1706.5085 37708.934 10000
```

### Question 5

To speed this up even further, we can write our own Metropolis-Hastings algorithm in Rcpp too and observe once again the same general results (note that the results won't be identical because of the stochasticity of simulations inside `synllk`).


```
cat mh.cpp
```

```
## #include <Rcpp.h>
## using namespace Rcpp;
##
## // [[Rcpp::export]]
## NumericVector metrop_Rcpp(Function lik, double init, int n, double scale){
##   double proposal;
##
##   NumericVector samples(n);
##   NumericVector alpha(n);
##
##   samples[0] = init;
```

```
##
##   for(int i=1; i<n; i++){
##       proposal = R::rnorm(samples[i-1], sqrt(scale));
##
##       alpha[i] = as<double>(lik(proposal))/as<double>(lik(samples[i-1]));
##
##       if(R::runif(0,1) < alpha[i]){
##           samples[i] = proposal;
##       }
##       else {
##           samples[i] = samples[i-1];
##       }
##   }
##
##   return samples;
## }

sourceCpp("mh.cpp")
samples = metrop_Rcpp(function(r) exp(synllk_Rcpp(r, 100, yobs)), log(40), 10000, 1)
plot(density(exp(samples)), main="r Approximate Posterior")
```

### r Approximate Posterior



N = 10000 Bandwidth = 0.5236

But now we have the benefit of a significant speed-up in obtaining these results: running everything in R is around 5 times slower than the full Rcpp implementation, which itself is about 40% faster than using `metrop` with an Rcpp implementation of `synllk`.

```
MH_R <- function() metrop(function(r) synllk(r, 100), initial=log(40), nbatch=50, scale=0.1)
MH_mix <- function() metrop(function(r) synllk_Rcpp(r, 100, yobs), initial=log(40), nbatch=50,
                             scale=0.1)
# Note below we exponentiate the likelihood function due to differences in the M-H implementation
MH_Rcpp <- function() metrop_Rcpp(function(r) exp(synllk_Rcpp(r, 100, yobs)), log(40), 50, 0.1)
microbenchmark(MH_R(), MH_mix(), MH_Rcpp(), times=100)
```

```
## Unit: milliseconds
##      expr      min      lq      mean      median      uq      max neval
##    MH_R() 115.46187 119.19204 121.40465 121.14991 122.51691 153.77629   100
##    MH_mix() 43.07378 46.44682 47.52050 47.53224 48.69843 53.08890   100
##    MH_Rcpp() 28.14810 28.76566 30.13895 29.88991 31.25612 33.43923   100
```

However, note that the above benchmark only ran the Metropolis-Hastings algorithms for 50 iterations. When we run this benchmark for a larger number of samples (below we use 500), the R implementation of `metrop` is faster than our Rcpp function `metrop_Rcpp`.

```
MH_R <- function() metrop(function(r) synllk(r, 100), initial=log(40), nbatch=500, scale=0.1)
MH_mix <- function() metrop(function(r) synllk_Rcpp(r, 100, yobs), initial=log(40), nbatch=500,
                             scale=0.1)
# Note below we exponentiate the likelihood function due to differences in the M-H implementation
MH_Rcpp <- function() metrop_Rcpp(function(r) exp(synllk_Rcpp(r, 100, yobs)), log(40), 500, 0.1)
microbenchmark(MH_R(), MH_mix(), MH_Rcpp(), times=5)
```

```
## Unit: milliseconds
##      expr      min      lq      mean      median      uq      max neval
##    MH_R() 908.4620 909.7799 915.6838 909.8278 910.7625 939.5870     5
##    MH_mix() 175.5941 176.2796 177.9067 177.3838 180.0604 180.2158     5
##    MH_Rcpp() 302.2229 302.8542 304.3112 303.4769 305.8246 307.1774     5
```